

A Numerical Formulation of Chan’s Hypervolume Algorithm

Cell masses, finite contractions, and preservation
of the $d/3$ complexity exponent

Technical report accompanying **FastHVChan**

5 October 2026

Abstract

Hypervolume measures the region dominated by a finite approximation set and is a central indicator in evolutionary multiobjective optimization. Its geometric formulation is the union measure of boxes sharing an anchor. We give a self-contained numerical formulation of Chan’s orthant recursion, written for both optimization and computational geometry readers. The integrals are evaluated by exact finite cell masses, prefix sums and monotone integer-cutoff constraints. No quadrature or general symbolic integration is required. The simplification is representational; Chan’s own step-function integrals also admit finite arithmetic evaluation.

The formulation starts directly from ordinary hypervolume under Lebesgue measure: coordinate weights are interval lengths, while pair constraints retain geometric dependence. We prove closure under numerical elimination and block-mass compression, then account for the entire recursive computation. A local checkpoint policy, optionally skipping compression when the state is already small, retains Chan’s $\tilde{O}_d(n^{d/3})$ arithmetic bound for every fixed $d \geq 4$. Thus the numerical treatment retains the exponent; it does not improve it. Dimensions two and three use $O(n \log n)$ sweeps.

The accompanying implementation covers dimensions two through ten. Experiments in 4D–6D use 64–1,024 points, with a 4,096-point 4D stress test, and compare the proved local policy, an earlier adaptive variant, and both Python Chan references. Completed timings, partial runs, and timeouts are kept distinct. Earlier small-input tests remain separate historical comparisons. The measurements do not establish an asymptotic exponent or universal practical superiority. An appendix develops the product-measure extension to dominated-set magnitude and its equivalence with HV; that extension is not needed by the main algorithm.

Keywords: hypervolume indicator; Klee’s measure problem; Lebesgue measure; product structure; orthants; numerical elimination; compression; evolutionary multiobjective optimization.

Contents

1	Purpose, audiences, and the precise claim	3
2	Dominated regions and ordinary hypervolume	4
3	Exact discretization and the numerical meaning of an integral	4
4	Why product structure helps, and what remains coupled	5
5	Finite contraction with monotone pair constraints	6
6	Compression is a mass-preserving change of index set	7
7	Geometric recursion and its invariant	8
8	A local compression policy that retains Chan’s bound	10
9	Software realization and validation of the local policy	13
10	Experiments with larger inputs in 4D–6D	15
11	Archived small-input experiments	18
12	Structural simplifications and limits	21
A	Dominated-set magnitude	22
B	A compact implementation and proof map	27

1 Purpose, audiences, and the precise claim

In multiobjective optimization, the output of a search is usually a set of mutually incomparable objective vectors. An indicator assigns a scalar value to this set, for assessment or for selection during the search. The hypervolume indicator rewards the size of the region dominated by the vectors relative to a reference point. Repeated indicator evaluations can therefore be a computational bottleneck; see the optimization background in [4].

For computational geometry, the same problem is a restricted union-measure problem. The boxes share an anchor, equivalently they are orthants clipped to a bounding box. This restriction is stronger than arbitrary Klee measure input. Chan’s general-box algorithm has an $O(n^{d/2})$ bound, whereas his orthant framework obtains $O(n^{d/3} \log^{O_d(1)} n)$ for fixed $d \geq 4$ [1]. The reference implementations in our experiments are named after these two exponents. A better exponent does not by itself imply a smaller runtime on a small input.

The objective here is to explain the orthant algorithm using numerical objects that can be inspected and compiled: ordered coordinate cells, their masses, monotone cutoff arrays, and signed sums of products. The presentation preserves the geometric recursion and the compression that make the bound possible. Simply replacing integrals by sums, or compiling a slow loop, would not be a sufficient complexity argument.

For the EMO reader. The main message is that the usual hypervolume indicator can be computed with interval lengths, prefix sums, and finite array operations, while retaining Chan’s asymptotic bound. The proof explains why the engine absorbs simple boxes before dividing the remaining instance and why compression is essential in high dimensions.

For the computational geometry reader. The new presentation specializes the functional algebra to a finite ordered grid of coordinate intervals. Its role is to make the implementation and its invariants explicit. We provide both a schedule that charges representation growth and a guarded version of that schedule. We do not claim a new orthant exponent, a new lower bound, or a new independent planar factorization.

Origin of the viewpoint. The magnitude viewpoint encouraged us to view the problem algebraically; the earlier [report in the HV4DMagnitude repository](#) documents that route [5]. In the end, the formulation can be developed directly for the common hypervolume indicator, because Lebesgue measure itself is a product measure. Magnitude is not required for the algorithm, its proof, or the measured HV timings. Appendix A contains the metric background, both indicator transformations, and the additional application of the same engine.

What the complexity statement covers. Our main bound applies to the local checkpoint algorithm of Sections 7–8. It includes the fixed local schedule already used by the standard-library numerical version and the new compiled local-checkpoint driver. The earlier experimental compiled driver uses an interval based on the *original* input size. We do not attach the local-schedule theorem to that different policy. Section 9 maps these statements to actual files.

2 Dominated regions and ordinary hypervolume

For a finite set $A \subseteq \mathbb{R}_{\geq 0}^d$, define

$$Q(a) = \prod_{i=1}^d [0, a_i], \quad D(A) = \bigcup_{a \in A} Q(a), \quad \text{HV}_d(A) = \lambda_d(D(A)).$$

The set $D(A)$ is a down-set in the nonnegative orthant: if $x \in D(A)$ and $0 \leq y \leq x$ coordinatewise, then $y \in D(A)$. We call a a generator or an upper corner. Repeated or dominated generators do not change the region. All statements allow tied and zero coordinates; no general-position hypothesis is needed.

For minimization vectors p and a reference vector r satisfying $p_i \leq r_i$, the conventional region is $\bigcup_p \prod_i [p_i, r_i]$. The reflection $x \mapsto r - x$ turns it into $D(\{r - p\})$. This reflection preserves volume, so the upper-corner convention computes the same hypervolume. The reference point still determines which region is measured.

For ordinary hypervolume, Lebesgue measure already has product structure:

$$\lambda_d = \bigotimes_{i=1}^d \lambda_1, \quad \lambda_d(Q(a)) = \prod_i a_i.$$

This is the algebraic starting point of the algorithm. A union of boxes need not be a Cartesian product; the geometric constraints must still be retained. No change of indicator is needed to obtain the product algebra.

3 Exact discretization and the numerical meaning of an integral

On axis i , sort the distinct positive generator coordinates and write $0 = c_{i,0} < c_{i,1} < \dots < c_{i,m_i}$. Use the cells

$$C_{i,0} = \{0\}, \quad C_{i,k} = (c_{i,k-1}, c_{i,k}] \quad (k \geq 1).$$

For Lebesgue measure their masses are simply

$$w_i(0) = 0, \quad w_i(k) = c_{i,k} - c_{i,k-1} \quad (k \geq 1).$$

The zero cell preserves the rank convention for zero coordinates; it contributes no ordinary volume. These interval lengths are the only initial weights required by the HV algorithm. Every generator becomes a rank vector $r(a)$. Membership in its box is exactly $k_i \leq r_i(a)$ on all axes. Therefore

$$\text{HV}_d(A) = \sum_{k_1=0}^{m_1} \cdots \sum_{k_d=0}^{m_d} \mathbf{1}[\exists a \in A : k \leq r(a)] \prod_i w_i(k_i). \quad (1)$$

This is an exact identity, not a mesh approximation. The algorithm stores only the separate axis arrays, of total length $O_d(n)$, and structured constraints. It does not construct the potentially $O(n^d)$ product grid displayed in the definition of the sum.

Consider a one-dimensional step density h and let $w(k)$ be its integral over cell k . For a union of whole consecutive cells,

$$\int_{\bigcup_{k=L}^U C_k} h(x) dx = \sum_{k=L}^U w(k) = P(U+1) - P(L), \quad P(t) = \sum_{k < t} w(k).$$

A step-function implementation stores breakpoints, density values, and an antiderivative represented by cumulative interval contributions. A cell-mass implementation stores the $w(k)$ and their prefixes directly. These are two finite arithmetic descriptions of the same operation.

The difference matters in implementation. Integer ranks handle strict and weak endpoints explicitly, and the bulk operations compile into loops over typed arrays. There is no general symbolic-expression manipulation and no quadrature error. Nevertheless, evaluating a known step-function antiderivative was already numerical arithmetic in the Python Chan reference. The simplification is an explicit finite representation, not the elimination of a computer algebra system that the reference never used.

For a concrete HV example, take $A = \{(2, 1), (1, 2)\}$. The axis coordinates are $(0, 1, 2)$ and the mass arrays are $(0, 1, 1)$. The only uncovered positive-area cell within the enclosing square is $(1, 2] \times (1, 2]$, indexed by $(2, 2)$. Thus the HV is $4 - 1 = 3$. This tiny grid is useful for explanation and testing; the high-dimensional solver avoids enumerating its analogue.

The finite algebra also accepts supplied masses from another product measure $\nu = \otimes_i \nu_i$: set $w_i(k) = \nu_i(C_{i,k})$ and replace $\text{HV}_d(A)$ on the left of (1) by $\nu(D(A))$. The coordinate cells still make each box-membership predicate constant. This extension changes the initial weights, not the recursive operations. Its cumulative-coordinate interpretation is given in Appendix A.

4 Why product structure helps, and what remains coupled

The product structure of Lebesgue measure lets the initial state be stored as d unary arrays instead of a table with one entry per product cell. It also lets us sum an unconstrained variable by a single scalar factor. Geometry enters through predicates. After absorbing a two-sided box or eliminating a variable, the state generally ceases to be one independent product. It remains a *sum of constrained products*, which is the useful closure property.

For example, in four dimensions let

$$A = \{(2, 1, 1, 1), (1, 1, 2, 1)\}.$$

Its HV is 3, but the product of the HVs of the projections on coordinate pairs $(1, 2)$ and $(3, 4)$ is $2 \cdot 2 = 4$. The projected calculation has lost the information about which extents occur together in a generator. Our representation keeps constraints connecting the pairs.

This also explains the scope of a “2 + 2 split”. In 4D we may eliminate axes $(1, 2)$ and then $(3, 4)$. Finite sums justify that grouping without an independence assumption. The implementation performs the two scalar eliminations within each pair; it does not provide a separate constant-time joint planar contraction. For odd d , the last coordinate is a singleton. No dummy dimension is needed.

Another benefit is that compression preserves masses directly. A block can contain positive-length cells and zero-mass endpoint cells. Its total mass is a sum. There is no need to divide by a cell's volume or reconstruct an artificial continuous density. The same finite operations also work for the supplied product masses described above.

5 Finite contraction with monotone pair constraints

Let p variables range over finite chains $E_i = \{0, \dots, s_i - 1\}$. The number p is at most d in a state and at most $2d$ during compression. A term has the form

$$T(k) = c \prod_i u_i(k_i) \prod_e \mathbf{1}[k_{j_e} \bowtie_e f_e(k_{i_e})], \quad i_e < j_e, \quad \bowtie_e \in \{\leq, \geq\}. \quad (2)$$

Coefficients and unary arrays may be signed. The integer cutoff arrays f_e are monotone. For each unordered variable pair we keep four possible classes: upper or lower inequality, each with a nondecreasing or nonincreasing cutoff. Same-class upper bounds combine by minimum and lower bounds by maximum. Cutoffs of opposite monotonicities stay separate. The sentinels -1 and s_j encode impossible upper and lower bounds.

A state F is a sum of S such terms. If $M = \sum_i s_i$, its storage is $O_p(SM)$; it is not a dense p -dimensional array. The signed representation is an algebraic device. It does not assert that the underlying initial measure has negative mass.

5.1 How comparisons stay in this class

Bounds on a variable are constants or monotone arrays in one other variable. Consider a comparison $A(k_i) \leq B(k_j)$. If $i = j$, evaluate it entrywise and multiply the resulting Boolean array into u_i . This array need not be monotone. If one side is constant, the comparison is likewise a unary mask. Otherwise, search the monotone array associated with the larger-index variable. For a nondecreasing B , for example, $A(k_i) \leq B(k_j)$ becomes

$$k_j \geq \min\{t : B(t) \geq A(k_i)\},$$

with an impossible-bound sentinel when the set is empty. Binary searches handle plateaus and strict inequalities using the appropriate first or last admissible position. The result is a monotone cutoff array. Nonincreasing arrays reverse the inequality direction in the same way. All four classes are needed for closure under these comparisons.

5.2 First-winner elimination

To sum out a variable q , rewrite each incident predicate as a lower or upper bound on q , inverting its monotone cutoff if needed. Include the domain endpoints. The feasible interval is

$$L(k) = \max_a L_a(k), \quad U(k) = \min_b U_b(k).$$

Each candidate depends on at most one surviving variable. Give the candidates a fixed order. For every pair (a, b) , let G_{ab} assert that L_a is the *first* maximizing lower bound, U_b is the *first* minimizing upper bound, and $L_a \leq U_b$. For the lower winner, the comparisons with earlier candidates are strict and those with later candidates weak; the upper winner has the analogous rule. Thus ties do not create overlaps.

With $P_q(t) = \sum_{r < t} u_q(r)$, the part of the sum on this winner region is

$$G_{ab}(k) [P_q(U_b(k) + 1) - P_q(L_a(k))]. \quad (3)$$

All factors in (2) that do not involve q remain in place. Each prefix evaluation in (3) is constant or unary in the variable supplying its bound. If the two bounds have different source variables, split the difference into two signed terms; if they have the same source, one unary difference suffices. In particular, a signed prefix array is never assumed monotone and is never inverted as a cutoff.

Lemma 5.1 (Numerical closure). *Eliminating one of p variables from one term produces at most*

$$B_p = 2(2p - 1)^2$$

terms of the form (2). For fixed p , it costs $O_p(M \log(M + 2))$ arithmetic and comparison operations per incoming term, apart from optional merging of equal resulting terms.

Proof. Each other variable contributes at most two lower and two upper candidates; with a domain endpoint there are at most $2p - 1$ of each. There are therefore at most $(2p - 1)^2$ winner pairs. Their guards are comparisons of the forms already described, hence can be represented by unary or pair masks. The winner regions partition the feasible assignments, and each produces at most two terms by (3). A dimension-bounded number of array scans and binary searches establishes the cost. No new variable or larger coordinate domain is introduced. $\square$

For a numerical illustration, take $u = (1, 2, 3, 4)$, so $P = (0, 1, 3, 6, 10)$. If the winning bounds are $L = 1$ and $U = 2$, the contracted mass is $P(3) - P(1) = 6 - 1 = 5$. Bounds that depend on other variables select different entries of this same prefix array. The winner guards explain how the solver shares that calculation while retaining the conditions under which those entries apply.

Equal terms may be normalized and merged, and zero terms discarded. These operations are valuable in practice, but the proof of closure and the global bound do not assume that cancellation or merging is successful. This distinction is essential: a small experimental term count cannot substitute for a worst-case argument.

6 Compression is a mass-preserving change of index set

Suppose m hard boxes remain in a recursive cell. On every axis, put a block boundary immediately after each hard-box cutoff below the cell's upper endpoint. The resulting consecutive blocks are $I_i(t_i) = [\ell_i(t_i), r_i(t_i)]$, with at most $m + 1$ blocks per axis. Define a coarse state by

$$(\text{Push } F)(t) = \sum_{k_i \in I_i(t_i) \forall i} F(k). \quad (4)$$

This is a specification by a finite sum, not an instruction to enumerate every product of blocks.

For each term, introduce the d new variables t_i with unary weight one and the comparisons $\ell_i(t_i) \leq k_i \leq r_i(t_i)$. There are now $2d$ variables. Eliminate just the old k_i variables

using Lemma 5.1. The result is a sum of terms on the block indices. The old variables can be eliminated in consecutive pairs, with a final singleton for odd d .

Lemma 6.1 (Compression and future masks). *If F has S terms and total axis length M , the compression emits at most $C_d S$ terms, where*

$$C_d = \prod_{p=d+1}^{2d} B_p.$$

Its work is $\tilde{O}_d(S(M + m + 1))$ before optional merging. The output has total axis length $O_d(m + 1)$ and storage $O_d(C_d S(m + 1))$. For every mask g constant on each product of blocks,

$$\sum_k F(k)g(k) = \sum_t (\text{Push } F)(t)g(t). \quad (5)$$

Proof. Apply Lemma 5.1 to the d old variables. The combined axis length is $M + O_d(m + 1)$ and decreases during elimination. Every surviving unary and cutoff array is indexed by a single coarse axis. For (5), partition the sum over the disjoint products of blocks and move the constant value of g outside each inner sum. $\square$

All surviving hard-box boundaries were used when making the blocks. Future restrictions and cuts based on these boxes are therefore aligned with them. Descendants introduce no new original coordinates: clipping uses current cell boundaries and surviving hard cutoffs. The compression can consequently be reused throughout the descendant recursion until another compression is appropriate.

Compression coarsens the coordinate vocabulary. It need not decrease the number of terms or the total stored entries immediately. The factor C_d is dimension dependent and can be large. In the actual natural compression probe, one compression changes one term into 296 terms. The benefit required by the proof is that the array lengths now depend on the remaining hard boxes rather than the old coordinate history. The proof must charge this growth in term count as well as the shorter arrays.

7 Geometric recursion and its invariant

At a node, let $\Gamma = \prod_i [l_i, h_i]$ be an inclusive integer-index cell, B the remaining boxes, and F the numerical state. The quantity computed is uncovered mass:

$$U(F, B, \Gamma) = \sum_{k \in \Gamma} F(k) \prod_{b \in B} (1 - \mathbf{1}[k \leq b]). \quad (6)$$

Initially $F(k) = \prod_i w_i(k_i)$. The final answer is the mass of the enclosing box minus U . After compression, F denotes block masses; the same invariant still applies.

7.1 Simplify before cutting

A box is disjoint if $b_i < l_i$ for some i and can be discarded. Otherwise, call axis i active if $b_i < h_i$. There are four cases.

1. With no active axis, the box covers the cell and $U = 0$.

2. With one active axis i , its covered region is a slab. Its complement tightens the lower bound to $b_i + 1$, or equivalently multiplies the state by the unary mask $\mathbf{1}[k_i > b_i]$.
3. For boxes with exactly two active axes $i < j$, combine the group using a suffix maximum:

$$f(k_i) = \max\{b_j : b \text{ in the group, } b_i \geq k_i\}, \quad \max \emptyset = -1.$$

The group is uncovered precisely when $k_j \geq f(k_i) + 1$. This is one nonincreasing lower cutoff, which can be absorbed into every term.

4. Every remaining hard box has at least three active axes.

Thus simple boxes are replaced by unary or pair constraints. Absorbing them does not multiply the number of terms. Shared extents produced by clipping are included automatically; they need not exist in the original input.

The compiled driver also crops the state to the current cell, removes constraints that become impossible or tautological on nonzero support, and merges equal states. These are exact restrictions, not approximations. All cell endpoints are inclusive. A split after rank c gives $[l, c]$ and $[c + 1, h]$; a slab covering through rank c leaves ranks starting at $c + 1$. Cropping shifts cutoff values by the target-axis offset. It does not recompute or insert initial cell masses.

7.2 Why three active axes produce the exponent $d/3$

Cycle through coordinate axes. At the current axis, let the cyclic positions of the axes be $0, 1, \dots, d - 1$. For each hard box and each triple J of its active axes, assign weight

$$\omega(J) = 2^{\sum_{j \in J} \text{pos}(j)/d}.$$

Let Φ be their total weight. For fixed d and nonempty hard input, $\Phi = \Theta_d(m)$, where m is the number of hard boxes. Choose a weighted median of the current-axis cutoff coordinates, using weights of the triples containing that axis, and split after it. Then advance the cyclic axis. If there are no candidate triples, simply advance the axis without a cut.

Lemma 7.1 (Weighted contraction). *With the exact weights above, each child after a cut and cyclic advance has potential at most $2^{-3/d}\Phi$. An advance with no candidate has the same bound and only one child.*

Proof. For a triple not using the current axis, the cyclic advance subtracts three from the sum of its positions, multiplying its weight by $2^{-3/d}$. A triple using the axis instead increases its position sum by $d - 3$. At most half the current-axis triple weight survives in either child: in the left child, a cutoff at or beyond the cut becomes inactive; in the right child, a cutoff at or below the cut is disjoint. Tied cutoffs at the median therefore cause no ambiguity. The combined factor for these triples is at most $\frac{1}{2}2^{(d-3)/d} = 2^{-3/d}$. Restriction and further absorption introduce no new active triples. When the axis has no candidate, only the first type of triple is present. $\square$

One full round of d advances has at most 2^d descendants, each with potential at most $\Phi/8$. Balancing 2^d branches against a size reduction by 8 suggests exponent $d/3$. This geometric observation alone does not bound the work: large numerical histories could still be carried into all those descendants. Section 8 supplies the missing accounting.

7.3 Integer median weights

Let n_0 be the original input size and $A_0 = (n_0 + 2)^2$. Instead of storing irrational weights, use

$$W_s = \lceil (A_0^d 2^s)^{1/d} \rceil, \quad 3 \leq s \leq 3d - 6. \quad (7)$$

Integer power comparisons compute these ceiling roots. They satisfy

$$2^{s/d} \leq W_s/A_0 \leq (1 + A_0^{-1})2^{s/d}.$$

Using their weighted median changes the potential contraction by at most $1 + A_0^{-1}$ per advance. Over $O_d(\log(n_0 + 2))$ advances this has bounded cumulative effect. For each fixed d , sufficiently large n_0 ensures $2^{-3/d}(1 + A_0^{-1})$ is bounded below one by a dimension-dependent constant. Smaller inputs lie below a dimension-dependent constant-size threshold, where the finite grid gives a bounded number of cuts and axis skips. The integer weights have $O_d(\log(n_0 + 2))$ bits. Their construction and use add only polylogarithmic work in our fixed-dimensional arithmetic analysis.

7.4 Small bases and low dimensions

If at most a fixed number b of hard boxes remain (the code uses $b = 2$), expand the product in (6). Its 2^b terms are signed rectangle sums, each evaluated by numerical elimination. This is constant-size inclusion–exclusion; it is not an exponential expansion over all input boxes.

In 2D, a descending first-coordinate sweep needs only the maximum active second-coordinate prefix mass. In 3D, the same sweep maintains the mass of a union of anchored rectangles. If $h(y)$ is the active maximum third-coordinate prefix mass, it is nonincreasing in the second rank. An inserted rectangle updates a prefix by $h(y) \leftarrow \max(h(y), v)$. The affected part is one interval: from the first height below v to the rectangle’s last second-coordinate rank. A lazy segment tree locates that first height and assigns the interval in $O(\log n)$, maintaining $\sum_y w_2(y)h(y)$. Sorting and all insertions take $O(n \log n)$ and $O(n)$ space. The same argument applies to supplied nonnegative initial masses.

8 A local compression policy that retains Chan’s bound

8.1 The checkpoint rule

A *block* is a consecutive group of geometric advances along a path. At a block entry, after absorption, let P be the total integer triple weight from (7), and put

$$Q = \lceil P/A_0 \rceil, \quad L = d \max \left\{ 1, \left\lceil \frac{\log_2 Q}{4d} \right\rceil \right\}. \quad (8)$$

Give the block a budget of L advances. Both cuts and axis skips consume one advance. At an uncompleted branch where the budget expires, simplify the node, then either:

1. compress onto its surviving hard-box blocks; or

2. skip that compression only if the total stored array entries plus term count are at most $K_d(m + 1)$, for a fixed dimension-dependent constant.

In *both* cases compute a fresh budget from the current Q . The implementation uses the equivalent threshold `compress_factor` $d^2(m + 2)$ with a fixed finite nonnegative factor (default eight). Taking factor zero forces compression at each checkpoint; it does not change the local checkpoint positions.

The skip is justified by a certificate about current storage. It is not an arbitrary postponement. Resetting the budget even after a skip is what keeps the next checkpoint local to the smaller subproblem.

Listing 1: Local-checkpoint recursion; all cell bounds are inclusive.

```
Solve(B, F, cell, axis, remaining):
    discard disjoint boxes; tighten slabs; crop F
    absorb two-sided groups as monotone pair masks
    if a box covers the cell or F is zero: return 0
    if at most b hard boxes remain: return Base(B, F, cell)

    if remaining == 0:
        if stored_size(F) > fixed_factor*d*d*(len(B)+2):
            F = Push(F, blocks at hard-box cutoffs)
            reindex B and cell onto the block axes
            remaining = unset      # reset even after a skip
    if remaining is unset:
        remaining = local_budget(current weighted potential)

    if the current axis has no eligible hard cutoff:
        return Solve(B, F, cell, next_axis, remaining-1)
    c = weighted_median(current-axis hard cutoffs)
    return Solve(left of c, next_axis, remaining-1)
        + Solve(right of c, next_axis, remaining-1)
```

Theorem 8.1 (Numerical Chan bound with guarded compression). *For fixed $d \geq 4$ and fixed base size, the local-checkpoint algorithm computes the ordinary hypervolume of a union of n anchored boxes in*

$$O_d(n^{d/3} \log^{c_d}(n + 2)) = \tilde{O}_d(n^{d/3}) \quad (9)$$

arithmetic and comparison operations, for a constant c_d depending only on d . This holds with compression at every checkpoint, or with the certified skip rule above. Initial coordinate sorting and interval length construction are included. The same bound applies to a supplied product measure whose cell masses take $O_d(n \log n)$ work to obtain. For $d = 2, 3$, the sweeps give $O(n \log n)$.

Proof. Correctness. Discretization is exact by (1). Absorption preserves (6); inclusive splitting partitions the cell disjointly. The base case is finite inclusion–exclusion followed by exact contractions. Lemma 6.1 preserves all later block-aligned restrictions. Skipping a compression leaves the state unchanged. Induction on the recursion proves the returned mass is correct.

Representation at a block entry. Let $N = \Phi$ after absorption, m be the hard-box count, S the term count, and M the total length of the common axis domains. At every nonroot block entry,

$$m = \Theta_d(N), \quad M = O_d(N + 1).$$

After compression the second relation follows from the $m + 1$ blocks per axis. After a skip it follows from the storage certificate: a nonempty term contains a unary array for every live axis, so M is no greater than the certified total storage. Within a block, clipping and absorption do not lengthen domains or increase S . Compression multiplies the term count by at most C_d from Lemma 6.1.

One block. Since $Q = \Theta_d(N)$, for sufficiently large N the budget is $L = \frac{1}{4} \log_2 N + O_d(1)$. The block contains $O_d(N^{1/4})$ nodes and leaves. At a surviving child checkpoint, Lemma 7.1 and the bounded integer-weight perturbation give

$$N' \leq c_d N^\beta, \quad \beta = 1 - \frac{3}{4d} < 1.$$

Each node can be charged $\tilde{O}_d(S^2 N)$, including coordinate scans, cut selection, contraction at an early base, or compression at a leaf. The square is a conservative deterministic allowance for comparing and merging terms; omitting merging gives a linear dependence on S . Thus the whole block costs $\tilde{O}_d(S^2 N^{5/4})$. In particular, no claim about beneficial cancellation is used.

Number of blocks and term growth. Above a sufficiently large fixed threshold, the recurrence $N' \leq c_d N^\beta$ gives $G = O_d(\log \log(n + 2))$ block generations on a path. Below that threshold, only a constant number of geometric advances remains: the potential contracts by a fixed factor, and a hard box has a fixed positive minimum triple weight. Consequently,

$$S \leq C_d^G = \log^{O_d(1)}(n + 2)$$

along every path from the one-term initial state. Skipped checkpoints do not increase S and do not invalidate the bound on generations, because they too restart the local budget.

Summing the work. Up to factors depending only on d , at macro-depth g a subproblem size is bounded by $cN_0^{\beta^g}$, and the number of block entries is at most

$$a_d^g N_0^{\frac{1}{4} \sum_{j < g} \beta^j} = a_d^g N_0^{\frac{d}{3}(1 - \beta^g)}.$$

Constants raised to g , as well as S^2 , are polylogarithmic in n . The sum of block costs at that depth is therefore bounded, up to such factors, by

$$N_0^{\frac{d}{3}(1 - \beta^g) + \frac{5}{4}\beta^g} = N_0^{\frac{d}{3} - (\frac{d}{3} - \frac{5}{4})\beta^g} \leq N_0^{d/3},$$

since $5/4 < d/3$ for every $d \geq 4$. Summing over all macro-depths adds only a polylogarithmic factor.

The root and low dimensions. Many boxes can be absorbed at the root while its arrays still have length $O_d(n)$. Its first block costs at most $\tilde{O}_d(nN_0^{1/4}) \leq \tilde{O}_d(n^{5/4})$, with $S = 1$, or $\tilde{O}_d(n)$ if it ends immediately. Both fit (9). Subsequent block entries obey the invariant already proved. The 2D and 3D sweep analysis supplies their separate bounds; substituting $d = 3$ into this block-work calculation would not be justified. $\square$

8.2 Meaning of the retained exponent

The proof is a numerical realization of Chan’s simplification, elimination, and compression framework [1]. His functional closure and grid-preserving simplification appear in Lemmas 4.4 and 4.5; the orthant bound is in Section 4.2. Our local schedule makes the representation-growth charge explicit and permits the storage-certified skip. The finite prefix operations cause no extra polynomial factor in n . In particular, 4D remains $\tilde{O}(n^{4/3})$.

d	Elimination grouping	Arithmetic bound	Lift variables
2	numerical sweep	$O(n \log n)$	—
3	sweep and skyline tree	$O(n \log n)$	—
4	$2 + 2$	$\tilde{O}(n^{4/3})$	8
5	$2 + 2 + 1$	$\tilde{O}(n^{5/3})$	10
6	$2 + 2 + 2$	$\tilde{O}(n^2)$	12
7	$2 + 2 + 2 + 1$	$\tilde{O}(n^{7/3})$	14
8	$2 + 2 + 2 + 2$	$\tilde{O}(n^{8/3})$	16
9	$2 + 2 + 2 + 2 + 1$	$\tilde{O}(n^3)$	18
10	$2 + 2 + 2 + 2 + 2$	$\tilde{O}(n^{10/3})$	20

The theorem is fixed-dimensional: its hidden constants and logarithmic exponents may depend substantially on d . It is an arithmetic-operation bound, not a rational bit-complexity result or a floating-point stability theorem. It concerns total indicator evaluation, not all exclusive contributions, gradient evaluation, or online update sequences.

8.3 Why the older adaptive timings do not prove this theorem

The earlier compiled adaptive driver uses a level interval of roughly $0.1d \log_2 n_0$, based on the original input size, combined with a size threshold. This does not restart local blocks at every smaller subproblem or skipped checkpoint. Moreover, even a locally chosen $0.1d \log_2 N$ interval would give the crude block-work exponent $1 + 0.1d$, which is $1.4 > 4/3$ in 4D. The constant in (8) is deliberate: block work grows like $N^{5/4}$, strictly below the target exponent in all dimensions under discussion.

Accordingly, the statement “the numerical treatment retains Chan’s complexity” is proved here for an explicit numerical algorithm and its local policy. It is not inferred from a familiar algorithm name, a compilation setting, or timing plots for a different policy.

9 Software realization and validation of the local policy

The repository [6] provides three relevant variants in `NumericalChanHVND`. They share the finite-mass problem formulation but should not be conflated.

1. `numerical_chen.py` is the standard-library numerical implementation with locally chosen blocks and compression at each checkpoint. Integer and `Fraction` inputs can be evaluated with exact arithmetic. It is covered by the fixed-local-policy case of Theorem 8.1.

2. `numerical_chan_compiled.py` is the previously published compiled adaptive driver. It adds cell pruning and native contractions, with the global-interval policy discussed above. The archived performance comparisons measure this variant.
3. `numerical_chan_local_compiled.py` is the local-checkpoint compiled driver accompanying this report. It reuses the same native kernel, cell restriction and base routines, and changes the scheduling to Listing 1. Theorem 8.1 covers its algorithmic policy, with fixed base size and compression factor, under the stated arithmetic-operation model.

All three support every integer dimension from two through ten. The theory is for any fixed dimension; ten is an implementation and testing range, not a mathematical barrier.

9.1 What is compiled

The file `compiled_terms.py` stores coefficients, unary arrays, and pair constraints in typed native states. Numba compiles complete restriction, masking, variable elimination, state merging, base sums, and block compression. The geometric recursion remains readable Python. In 2D and 3D the numerical sweeps are compiled instead. The earlier hybrid Numba version compiled only smaller array helpers while Python retained much of the term construction and dispatch overhead.

The kernel does not enable fast-math or parallel execution. Coefficients and unary masses are binary64 values, cutoff indices are integers, and median weights are the separately constructed Python integers. Dictionary copies share arrays only while those arrays are immutable; modifications allocate replacement arrays. State hashes are checked against full contents before equal terms are merged. Hash collisions cannot justify an algebraic equality. A limit on collision-chain comparisons may leave equivalent terms unmerged, which preserves their sum and does not affect the raw-emission proof.

Compilation changes the cost of an operation, not the identity it evaluates. Reusing the native kernels under the local policy does not require a new elimination rule. Conversely, enabling Numba does not establish either exact real arithmetic or numerical stability. Signed cancellation can magnify floating-point errors even when the final indicator is positive.

9.2 Minimal use

From the repository root, in an environment with NumPy and Numba:

```
from NumericalChanHVND.numerical_chan_local_compiled import (
    hypervolume, NumericalChanLocalCompiled)

points = [(2, 1, 1, 1), (1, 2, 1, 1)]
assert abs(hypervolume(points) - 3.0) < 1e-10

solver = NumericalChanLocalCompiled(dimension=4)
value = solver.compute(points)
print(solver.stats)
```

For minimization data, replace each point by its nonnegative extent vector $r - p$ before calling the solver. The public validation rejects negative or nonfinite extents. An empty

input and a single all-zero generator both have HV zero. No dominance filter is required for correctness.

The local driver intentionally exposes no global interval or compression-disable option. Its compression factor is a fixed constant in the theorem; factor zero means compression at every scheduled checkpoint. Arbitrarily changing the block budget or making policy parameters depend on input size is outside the stated theorem.

9.3 Validation of the newly proved policy

The script `verify_local_compiled.py` and its saved JSON record check the new driver independently of the old timing experiment. They include 232 exact small-instance checks across dimensions 2–10, 24 additional exact checkpoint-recursion instances, and 3,538 observed cyclic-advance and budget invariants. Additional product-measure checks are described in Appendix A.4. The trace checks that a skipped axis consumes a step and that a checkpoint resets its budget after either decision.

A 64-point four-dimensional permuted antichain reaches two successive checkpoint generations. With factor zero it executes 63 compressions; with a sufficiently large fixed factor it executes 63 certified skips. Both settings agree with the independent Chan $d/2$ reference for HV. These deliberately chosen settings validate both branches of the policy; they are not performance recommendations.

The common-extent 1,552-box family described in Section 11 also executes six *default local-policy* compressions. Its HV is 6 372 330 876, agreeing with the independent Chan $d/2$ computation. These six events occur on different branches; this case alone does not demonstrate repeated compression on one path.

All checks pass with maximum scaled discrepancy 5.6844×10^{-13} , using relative tolerance 2×10^{-9} and absolute tolerance 2×10^{-10} . Cached import in a fresh process is also checked. The shared kernel’s separate suite tests signed states, all cutoff orientations, inclusive crops, repeated compression, and parent-state immutability. These tests corroborate the implementation; the complexity conclusion rests on Theorem 8.1, not on a finite test set.

10 Experiments with larger inputs in 4D–6D

This additional experiment measures the local-checkpoint implementation covered by Theorem 8.1, alongside the earlier global-interval compiled variant and both Python Chan references. The solvers were unchanged from source revision 095fa73; hashes identify the exact sources and inputs. Here n denotes input generators, one per anchored box.

The spherical inputs use $n = 64, 128, 256, 512, 1024$ in each of dimensions four, five and six, and $n = 4096$ in four dimensions. A second family on the positive simplex uses $n = 64, 256, 1024$ in each dimension. Smaller inputs are nested prefixes within each family and dimension. Spherical coordinates are absolute Gaussian samples plus 0.01, normalized in Euclidean norm; simplex coordinates are exponential samples plus 0.01, normalized by their sum. Both constructions produce positive nondominated fronts.

Each fresh sequential worker makes one first call and three warm calls, with a new solver on each call. Warm medians include input conversion and preprocessing but exclude imports and first use. Existing compilation caches are retained. The whole worker,

d	n	Local compiled	Earlier compiled	Chan $d/3$ Python	Chan $d/2$ Python
4	64	0.0518	0.0541	0.1243	0.0245
4	128	0.1431	0.1417	0.3472	0.0684
4	256	0.3427	0.3366	0.8913	0.1622
4	512	0.8485	0.9259	2.1206	0.3859
4	1024	3.4890	4.9178	4.7949	0.8810
4	4096	TO[0]	TO[0]	TO[0]	4.5595
5	64	0.2208	0.2145	0.6844	0.0878
5	128	0.6220	0.6149	1.9354	0.2568
5	256	1.9115	1.8885	6.0686	0.8202
5	512	5.2159	5.2239	TO[1]	2.0171
5	1024	TO[1]	TO[0]	TO[0]	5.0315
6	64	0.7754	0.7949	3.0187	0.2772
6	128	3.0644	3.0222	10.9978	1.2158
6	256	10.3680	10.2754	TO[0]	3.9232
6	512	TO[0]	TO[0]	TO[0]	TO[2]
6	1024	TO[0]	TO[0]	NR	TO[0]

Table 1: Sphere fronts at larger input sizes: warm median *seconds*. Local compiled uses the proved local-checkpoint policy; earlier compiled uses the global-interval policy. TO[k] means worker timeout with k complete warm calls; NR means not run after two smaller-size timeouts.

including imports and all four calls, has a 45-second budget. A timeout therefore does not imply that one call took 45 seconds. Available partial timings and values are archived. After two consecutive timeouts in a solver/family/dimension series, larger sizes in that series are explicitly marked not run; no runtime is inferred for them.

All default compression settings remain enabled. There is no dominance prefilter, fast-math or concurrent numerical worker. The processor is an AMD Ryzen 7 5800H with 16 logical CPUs. The environment is the same Windows/Python/Numba setup described in Section 11. Ordinary system load is uncontrolled.

Tables 1 and 2 report complete warm medians. Table 3 retains first-call observations on selected large inputs, including partial workers.

There are 25 inputs and 100 planned algorithm/input rows: 74 complete workers, 25 timeouts, and 1 not run by the escalation rule. The local solver completes 18 of its 25 rows. The tables retain all planned sizes so the missing large cases remain visible.

Against the earlier compiled solver, local compiled is faster on 8 of 18 complete pairs, with median reference/local ratio 0.991. These ratios omit censored pairs and are not universal speed guarantees.

Against Python Chan $d/3$, local compiled is faster on 16 of 16 complete pairs, with median reference/local ratio 2.88. These ratios omit censored pairs and are not universal speed guarantees.

Against Python Chan $d/2$, local compiled is faster on 0 of 18 complete pairs, with median reference/local ratio 0.417. These ratios omit censored pairs and are not universal speed guarantees.

d	n	Local compiled	Earlier compiled	Chan $d/3$ Python	Chan $d/2$ Python
4	64	0.0589	0.0543	0.1566	0.0262
4	256	0.3693	0.3574	0.9423	0.1621
4	1024	4.0905	5.4003	5.4955	0.9622
5	64	0.2168	0.2385	0.7958	0.1137
5	256	2.1954	2.2019	7.7507	0.9182
5	1024	TO[0]	TO[0]	TO[0]	6.1703
6	64	0.8562	0.8410	3.2904	0.3552
6	256	TO[2]	TO[2]	TO[0]	5.2442
6	1024	TO[0]	TO[0]	TO[0]	TO[0]

Table 2: Simplex fronts at larger input sizes: warm median *seconds*. Local compiled uses the proved local-checkpoint policy; earlier compiled uses the global-interval policy. TO[k] means worker timeout with k complete warm calls; NR means not run after two smaller-size timeouts.

Front	d	n	Local compiled	Earlier compiled	Chan $d/3$ Python	Chan $d/2$ Python
Sphere	4	1024	5.868	7.325	4.772	0.885
Sphere	4	4096	42.565	42.937	23.367	4.539
Sphere	5	1024	22.582	31.550	43.535	4.932
Sphere	6	512	36.826	36.728	—	12.678
Sphere	6	1024	—	—	NR	34.491
Simplex	4	1024	6.377	7.522	5.494	0.959
Simplex	5	1024	25.191	34.373	—	6.544
Simplex	6	1024	—	—	—	43.563

Table 3: Retained first-call seconds on the largest inputs, including workers that later timed out. These are single observations, not warm medians; compiled first calls can include kernel initialization and cache loading. Imports are excluded. A dash means no first value returned within the worker budget; NR means not run.

For four dimensions and $n = 4096$, local compiled returns its first value in 42.565 seconds, versus 23.367 seconds for Python Chan $d/3$. Neither completes all three warm calls. This first-call observation shows why the complete-pair speedups should not be extrapolated to larger inputs; it is not a warm-median comparison.

Compression is executed on 4 completed local inputs. Per-input checkpoint, compression and peak-term counts are retained with the raw data. The size rule is unchanged; absence of compression in an individual input is a recorded algorithm decision, not a disabled feature.

All 321 returned values in 91 rows are checked against an available Chan $d/2$ value; there are 0 mismatches and the maximum scaled error is 4.4409×10^{-16} . This checks every retained first and warm value, including partial workers. A completed first reference value remains valid even if that worker later exceeds its budget. These are cross-implementation comparisons; the larger sets were not evaluated by an exponential exact inclusion-exclusion oracle.

Among complete local workers, the ratio of slowest to fastest warm call has median 1.05 and maximum 1.26. The sample is one seeded front per dimension/family and three timed repetitions. It demonstrates performance and limits at larger n , not a fitted asymptotic exponent. The worst-case bound continues to rest on Theorem 8.1. Full inputs, statuses and retained samples are in `benchmarks/larger_4d_6d/results`.

11 Archived small-input experiments

These measurements concern the previously published *global-interval adaptive* implementation, archived with revision 38587e7 [6]. They do *not* measure the new local-checkpoint driver `numerical_chan_local_compiled.py`, which uses a different schedule. Source hashes identify the files actually timed.

11.1 Inputs, implementations, and measurement protocol

The main suite contains 45 standard inputs: for each dimension from 2 to 10, two spherical fronts, two tied grids, and one magnitude instance on the smaller spherical front. Three additional inputs force compression in dimensions 4, 7, and 10, giving 48 inputs altogether. Spherical corners are positive, normalized vectors formed from absolute Gaussian samples with a positive offset, and hence mutually nondominated. The grids can contain ties, duplicates, and dominated corners. Point counts decrease with dimension; this is a collection of small practical examples, not an experiment estimating an asymptotic exponent. The auxiliary indicator is defined in Appendix A; its presence in this archived suite does not enter the derivation of the HV algorithm.

Four implementations are compared: compiled numerical, original numerical Python, and the repository’s Python Chan $d/3$ and $d/2$ implementations. Chan $d/3$ is omitted in 2D. The compiled solver uses sweeps in 2D/3D and native contractions with Python geometric recursion in 4D–10D. Numerical Python retains its earlier local block schedule, so this compares complete implementations rather than compilation alone. Compression remains enabled. No dominance filtering, fast-math, or parallel execution is used. Solvers receive equivalent unions; magnitude uses the affine HV transformation for the Chan references.

d	n	Compiled numerical	Numerical Python	Chan $d/3$ Python	Chan $d/2$ Python
2	128	0.474	0.770	—	5.313
3	128	0.981	1.609	44.524	19.108
4	24	16.995	1877.982	28.632	6.216
5	16	16.971	1434.933	78.805	7.853
6	12	27.651	1634.712	54.249	9.206
7	8	20.201	759.785	66.189	6.179
8	10	130.842	11894.048	182.583	24.820
9	6	30.459	timeout	51.050	5.342
10	6	62.646	6737.424	96.593	11.206

Table 4: Larger spherical input in each dimension: warm median milliseconds for the pre-existing global-interval adaptive compiled implementation and three references. These measurements do not apply to the new local-checkpoint driver. Sizes vary with dimension.

Sequential fresh workers perform one separately recorded first call and three timed calls, each with a fresh solver. Warm medians include conversion and preprocessing but exclude imports; verification warmed the compilation caches. The 45-second worker budget includes startup, imports, and every call. A timeout is not a per-call lower bound. The recorded processor identifier is AMD64 Family 25 Model 80 Stepping 0, with 16 logical CPUs. The environment used Windows 11, Python 3.12.14, NumPy 2.5.3, and Numba 0.68.0; CPU affinity was not fixed. All raw samples, inputs, counters, signatures, and source hashes are archived [6].

Of 187 algorithm–input workers, 186 completed and one timed out: numerical Python on the six-point 9D sphere. All 48 compiled workers completed. Across ordinary-hypervolume inputs in dimensions 4–10, compiled numerical was faster than Chan $d/3$ on 26 of 28 pairs, with median ratio of reference time to compiled time 2.02. Against numerical Python, it won 26 of 27 completed pairs, with median ratio 6.29. Chan $d/2$ was faster on all 28 ordinary-hypervolume pairs in these dimensions. These completed-pair ratios exclude magnitude, forced compression, and timeouts. They provide no general claim that the compiled solver is the fastest hypervolume implementation.

Variability is material. Across the 48 compiled workers, the ratio of slowest to fastest warm repetition had median 1.28 and maximum 2.74; five workers exceeded a ratio of two. There were no deliberately concurrent solver or oracle jobs, but ordinary system load was uncontrolled. Three samples do not provide reliable confidence intervals, and small differences should not be overinterpreted.

11.2 What the compression experiments establish

All 35 standard inputs in dimensions 4–10 executed *zero* compressions in the compiled solver. Their adaptive size gates were checked 994 times, but the states stayed below the threshold. Thus the main timing table measures the benefit of pruning and native evaluation on these inputs; it does not measure the cost of required compression. The three forced cases use the first eight corners of a 4D spherical input, padded with unit extents for 7D and 10D. With interval two and threshold factor zero, each executes one

Event	Hard boxes	Size before	Threshold	Terms after	Size after
1	6	1710	1024	296	6936
2	3	1108	640	111	1665
3	4	1205	768	130	2237
4	4	1100	768	105	1785
5	4	1277	768	144	2468
6	4	974	768	110	1870

Table 5: All default compression events in the separate 1,552-box probe. Each starts with one term. Size counts stored scalar array entries plus terms. Events occur in different recursion branches: the maximum number of compressions along a path is one.

compression and reaches a peak of 135 terms. All values agree with the reference. These checks exercise the implementation, but are excluded from the performance aggregates.

A separate structural input exercises the global-interval implementation’s *unchanged defaults*. For each of the six coordinate pairs in 4D, it contains 256 corners with pair extents $(t, 257 - t)$, $1 \leq t \leq 256$, and extent 300 in the other coordinates. Add 16 corners $(161 + i, 176 - i, 161 + j, 176 - j)$ with $0 \leq i < 16$ and $j = 3i \bmod 16$. This gives 1,552 boxes. The first 1,536 become pair constraints, leaving long arrays with few remaining hard boxes. Shared extents are intentional; this is not a general-position performance sample.

The compiled run performs six compressions and returns 6 372 330 876, exactly the floating-point value returned by both Chan references. Table 5 records all six events. Every event starts above the size threshold, yet every resulting state is larger by the stored-entry count. In particular, the first compression changes one term into 296 terms. Compression preserves the required cell masses and coarsens the coordinate domains; it need not immediately reduce the complete representation.

This probe visited 47 nodes, emitted 6,358 raw terms, and performed 1,352 variable eliminations. Its single first-use computations took 2.432 seconds for compiled numerical, 1.235 seconds for Chan $d/3$, and 0.459 seconds for Chan $d/2$, using existing caches. They are *not warm medians* and are excluded from Table 4 and all timing ratios. Chan $d/3$ performed no compression: its breakpoint representation and the numerical array representation have different size measures. This experiment establishes that default compression is operational, not that it improves runtime or storage. Repeated compression along one path is checked in the separate correctness suite, not established by this probe.

11.3 Correctness evidence and its limits

Every available first-call and warm value was compared with Chan $d/2$. An independent oracle also evaluated exact rational inclusion–exclusion on the binary64 input coordinates for the 34 saved inputs with at most 12 points, including all three forced cases. All 543 recorded values passed; the largest scaled error was 6.3811×10^{-16} , where scaling divides absolute error by $\max(1, |\text{reference}|)$. The comparison tolerances were 2×10^{-9} relative and 2×10^{-10} absolute. Oracle input hashes match the archived dataset and timing files.

The proofs concern exact arithmetic; the compiled implementation uses binary64 and signed intermediate sums, without a general floating-point error bound. Local algebraic closure also does not establish the old adaptive driver’s global runtime. The

local-checkpoint analysis is separate from this archived benchmark. Section 10 now measures that policy on larger inputs; the earlier speedup ratios are not applied to it.

12 Structural simplifications and limits

The structural simplification can be stated precisely. Initially a product measure is stored as one unary factor per axis. Local geometric simplification replaces one- and two-sided boxes by unary and monotone pair predicates. Summing one variable introduces only a fixed-dimensional number of terms of the same form. Compression replaces long coordinate histories by masses on the surviving hard-box grid. The local schedule limits the number of representation-expanding checkpoints to $O_d(\log \log n)$ on a path. Together these properties preserve the $d/3$ exponent.

For ordinary HV, all initial masses are coordinate interval lengths. The algebraic formulation, including its product factors and finite prefix operations, therefore follows directly from Lebesgue measure. The extension in Appendix A is additional functionality; the benchmarked HV algorithm needs no metric kernel or indicator conversion.

There are also clear limits. Multiplying answers from coordinate-pair projections loses dependence. After masks are introduced, independent product factors are insufficient; the pair constraints and signed sums are essential. Compression can create many such terms before any merging occurs. The finite closure constants and the logarithmic exponent in the proof may be large. These facts are consistent with a theoretically better exponent losing to a simpler method on the tested inputs.

The present result leaves the stricter target of an unrestricted $\tilde{O}(n^{4/3-\varepsilon})$ 4D algorithm open. Obtaining it would require a stronger global bound or an additional structural reduction. Neither replacing integrals by finite sums nor grouping coordinates in pairs supplies that step. No new conditional lower bound is established here.

For EMO software, this gives a numerical implementation of the standard hypervolume indicator. For geometric analysis, it gives an explicit finite representation and an accounting of retained histories. The numerical formulation retains Chan’s exponent under the proved compression policy.

A Dominated-set magnitude

This appendix records an additional application of the numerical algebra. The development of the HV algorithm and its complexity proof in the main text requires only Lebesgue measure. The magnitude viewpoint encouraged the algebraic formulation in the earlier research report [5]; it is retained here as motivation and as an extension, not as a prerequisite for the ordinary hypervolume algorithm or as a source of a faster exponent.

A.1 Cumulative-coordinate reduction for product measures

Let ν_i be a nonnegative measure on the i th coordinate axis, finite on the bounding interval, and let

$$\nu = \bigotimes_{i=1}^d \nu_i, \quad F_i(t) = \nu_i([0, t]).$$

Product structure means that a box has mass $\nu(Q(a)) = \prod_i F_i(a_i)$. It does *not* mean that a union of boxes is a Cartesian product. The following elementary reduction captures exactly what can be factored.

Proposition A.1 (Cumulative-coordinate reduction). *For finite $A \subseteq \mathbb{R}_{\geq 0}^d$, put $F(A) = \{(F_1(a_1), \dots, F_d(a_d)) : a \in A\}$. Then*

$$\nu(D(A)) = \text{HV}_d(F(A)).$$

If all required cumulative values are supplied, the reduction needs $O(nd)$ arithmetic and coordinate operations, preserves dimension, and does not increase the number of generators.

Proof. Every nonempty intersection of input boxes is $Q(m_J)$, where $m_{J,i} = \min_{a \in J} a_i$. Since each F_i is nondecreasing, $F_i(m_{J,i}) = \min_{a \in J} F_i(a_i)$. The mass of the intersection therefore equals the volume of the intersection of its transformed anchored boxes. Finite inclusion–exclusion proves the equality for the unions. The empty input gives zero on both sides. $\square$

The proof is an identity, not an exponential-time implementation: inclusion–exclusion is used to establish the reduction; a union algorithm computes its right side. Atoms and flat parts of the cumulative functions are allowed. This is generally a transformation of *generators*, not an ordinary invertible change of coordinates on the entire region. When cumulative values require preprocessing, their construction cost must be added to the HV computation.

A.2 What dominated-set magnitude means

A.2.1 From similarity to metric magnitude

For a finite metric space $X = \{x_1, \dots, x_m\}$, form the similarity matrix $Z_{ij} = e^{-\text{dist}(x_i, x_j)}$. When Z is invertible, its magnitude is $\mathbf{1}^T Z^{-1} \mathbf{1}$. Equivalently, solve $Zw = \mathbf{1}$ and sum the weights. Closely spaced points have high similarity and do not behave like entirely independent units. This is the metric analogue of counting developed in the magnitude literature [3].

For a compact region, an applicable continuous formulation uses a finite signed *weight measure* W supported on the region, with

$$\int_D e^{-\text{dist}(x,y)} dW(y) = 1 \quad \text{for every } x \in D. \quad (10)$$

In the positive-definite setting used here, existence of such a weight measure gives $\text{Mag}(D) = W(D)$ [3]. We construct one explicitly for anchored box unions, so no limiting discretization of the metric space is needed for their computation. The quantity of interest is $\text{Mag}(D(A))$, *not* the magnitude of the finite generator set A .

For $\text{dist}(x, y) = \|x - y\|_1$, the key identity is

$$e^{-\|x-y\|_1} = \prod_{i=1}^d e^{-|x_i - y_i|}. \quad (11)$$

The sum in the distance becomes a product in the kernel. This tensor structure connects metric magnitude with the finite product algebra of the main text. The resulting algorithm evaluates no exponentials or similarity-matrix inverses.

For completeness, positivity and the weight-measure rule have a short energy explanation here. The one-dimensional Laplace kernel has positive Fourier transform $2/(1 + \xi^2)$; its tensor product is positive definite. For finite signed measures η on D , write $E(\eta) = \iint e^{-\|x-y\|_1} d\eta(x) d\eta(y)$. The energy characterization of magnitude is $\sup_{\eta \neq 0} \eta(D)^2 / E(\eta)$ in this setting. If W satisfies (10), then $\langle \eta, W \rangle = \eta(D)$ and $E(W) = W(D)$. Cauchy-Schwarz gives $\eta(D)^2 / E(\eta) \leq W(D)$, with equality for $\eta = W$. This justifies the rule used in the proof below and explains its relation to the finite matrix formula.

A.2.2 A direct proof for anchored boxes

For $a \geq 0$, let

$$w_a = \frac{1}{2}(\delta_0 + \delta_a + \lambda|_{[0,a]}).$$

When $a = 0$, the coincident endpoint atoms give $w_0 = \delta_0$. Splitting an elementary integral at x shows that, for every $x \geq 0$,

$$\int_{[0,a]} e^{-|x-y|} dw_a(y) = e^{-(x-a)_+}, \quad w_a([0,a]) = 1 + a/2, \quad (12)$$

where $t_+ = \max(t, 0)$. For $0 \leq x \leq a$, the integral of the density is $1 - (e^{-x} + e^{-(a-x)})/2$; the two atoms supply the missing terms. For $x > a$, all contributions have the common factor $e^{-(x-a)}$.

Tensor products of these interval measures give a weight measure for $Q(a)$ and yield $\text{Mag}(Q(a)) = \prod_i (1 + a_i/2)$. More is true.

Theorem A.2 (Product computing measure for dominated sets). *Let A be any finite set of nonnegative corners. With*

$$\mu = \bigotimes_{i=1}^d (\delta_0 + \tfrac{1}{2}\lambda),$$

the ℓ_1 magnitude of its dominated region satisfies

$$\text{Mag}(D(A)) = \mu(D(A)). \quad (13)$$

The magnitude of the empty region is taken to be zero.

Proof. Assume $A = \{a^1, \dots, a^n\}$ is nonempty. For nonempty $J \subseteq \{1, \dots, n\}$, let m_J be the coordinatewise minimum of its corners and let $W_J = \otimes_i w_{m_{J,i}}$. Define

$$W = \sum_{\emptyset \neq J \subseteq [n]} (-1)^{|J|+1} W_J.$$

This is a finite signed measure supported on $D(A)$. Its potential at a nonnegative point x is the same signed sum of $\exp(-\sum_i (x_i - m_{J,i})_+)$, by (11)–(12). If $x \in D(A)$, choose k with $x \leq a^k$. Pair every nonempty J not containing k with $J \cup \{k\}$. Their potentials agree: replacing $m_{J,i}$ by $\min(m_{J,i}, a_i^k)$ leaves $(x_i - m_{J,i})_+$ unchanged because $a_i^k \geq x_i$. All these pairs cancel. The unpaired subset $\{k\}$ has potential one. Thus W satisfies (10).

Its total mass is

$$W(D(A)) = \sum_{\emptyset \neq J \subseteq [n]} (-1)^{|J|+1} \prod_i (1 + m_{J,i}/2).$$

The last product is also $\mu(Q(m_J))$. Inclusion–exclusion for the positive measure μ now gives $W(D(A)) = \mu(D(A))$. $\square$

Equation (13) is the dominated-set formula in [2]. The proof above spells out why it is legitimate on this class, without assuming that magnitude is a valuation on arbitrary sets. The positive *computing measure* μ is not in general the kernel weight measure W . For instance, the translated interval $[1, 2]$ has metric magnitude $3/2$ by translation invariance, but its mass under the fixed-anchor measure $\delta_0 + \lambda/2$ is $1/2$. Consequently, this formula does not directly identify arbitrary Klee-measure instances with the metric magnitude of their unions.

A.2.3 Faces, Pareto compliance, and scale

Expanding the product measure in (13) gives

$$\text{Mag}(D(A)) = \sum_{I \subseteq [d]} 2^{-|I|} \lambda_{|I|}(\pi_I D(A)). \quad (14)$$

For a nonempty region the zero-dimensional term is one; for the empty region all terms are zero. The reason projections appear is specific to down-sets: the section obtained by setting coordinates outside I to zero coincides with the projection onto I . In two dimensions the formula is

$$1 + \frac{1}{2}(\text{horizontal extent} + \text{vertical extent}) + \frac{1}{4}(\text{area}).$$

In higher dimensions it combines all coordinate-projection volumes, including the full-dimensional HV. The algorithm need not enumerate these 2^d projections; one product measure contains them all.

Proper inclusion of two finite dominated regions strictly increases μ , hence magnitude. To see strictness, take $p \in D(B) \setminus D(A)$. Closedness of $D(A)$ leaves a sufficiently small lower rectangle near p outside $D(A)$; use intervals on the positive coordinates of p and origin atoms on its zero coordinates. This rectangle has positive μ mass and lies in $D(B)$. If $p = 0$, its atomic mass is already positive. This explains Pareto compliance at the region level. Adding an already dominated generator changes neither indicator. A new

box of zero full-dimensional volume can change magnitude through a lower-dimensional face even when it leaves ordinary HV unchanged.

Magnitude uses a metric and therefore a length scale. For the distance $\tau\|x - y\|_1$, with $\tau > 0$, the computing measure becomes $\otimes_i(\delta_0 + \frac{\tau}{2}\lambda)$. Its value is a polynomial in τ , whose degree- d term is $(\tau/2)^d \text{HV}_d(A)$ (and vanishes when the full-dimensional HV is zero). Changing objective units without correspondingly specifying the metric changes the indicator. Product structure gives a clear algebraic description of this choice; it does not remove the choice.

A.3 Magnitude and hypervolume: both reductions

Corollary A.3 (Magnitude to HV). *For finite nonnegative corners,*

$$\text{Mag}(D(A)) = \text{HV}_d(\{\mathbf{1} + a/2 : a \in A\}). \quad (15)$$

Proof. Apply Proposition A.1 with $F_i(t) = 1 + t/2$ and Theorem A.2. $\square$

The anchor on the right of (15) remains zero. The affine image of $[0, a_i]$ would instead be $[1, 1 + a_i/2]$. Thus this is not a statement that translating a region adds volume. Transforming the generators and rebuilding their anchored boxes includes precisely the face contributions represented by the atoms.

Proposition A.4 (HV to magnitude). *For an HV input $B \subseteq \mathbb{R}_{\geq 0}^d$, discard corners with any zero coordinate. If none remain, the HV is zero. Otherwise put*

$$\delta_i = \min_{b \in B} b_i > 0, \quad c_i(b) = 2(b_i/\delta_i - 1).$$

Then

$$\text{HV}_d(B) = \left(\prod_i \delta_i \right) \text{Mag}(D(\{c(b) : b \in B\})). \quad (16)$$

Proof. The discarded boxes have zero d -dimensional volume. Scaling the remaining i th coordinates by $1/\delta_i$ multiplies their HV by $\prod_i \delta_i^{-1}$ and gives extents at least one. For these normalized corners, $1 + c_i(b)/2 = b_i/\delta_i$. Apply (15) and undo the volume scale. $\square$

Both directions require one indicator evaluation and $O(nd)$ preprocessing. They preserve dimension and input size. Therefore an arithmetic upper bound transfers in both directions, and so does a conditional lower bound within a computational model supporting these arithmetic transformations. This observation alone proves no new hardness result. Bit complexity and restrictions on integer coordinates require their own accounting. Zero coordinates created by (16) must be retained in the *magnitude* instance. Discarding them would invalidate the reduction.

Example A.5 (One example in both directions). Take $A = \{(2, 1), (1, 2)\}$. Its HV is $2 + 2 - 1 = 3$, whereas its magnitude is

$$3 + 3 - \frac{9}{4} = \frac{15}{4} = 1 + \frac{1}{2}(2 + 2) + \frac{1}{4} \cdot 3.$$

The forward reduction gives corners $(2, 3/2)$ and $(3/2, 2)$, with HV $15/4$. The reverse reduction applied to the original HV instance gives $(2, 0)$ and $(0, 2)$: the magnitude of these two anchored line segments is $2 + 2 - 1 = 3$. Adding $(3, 0)$ to the original set leaves its HV at 3 and increases its magnitude to $17/4$.




Figure 1: The transformed generators define a new anchored region. The left region’s magnitude equals the right region’s ordinary area. The origin atom and lower-dimensional contributions are represented by the enlargement of anchored extents, not by a translation of the shaded set.

A.4 Using the same numerical engine

On the cells of Section 3, the initial masses are

$$w_i(0) = 1, \quad w_i(k) = \frac{c_{i,k} - c_{i,k-1}}{2} \quad (k \geq 1).$$

Replacing the Lebesgue masses by these masses changes neither the cutoff algebra nor the recursion. Prefix sums include the original origin atom. Cropping shifts rank indices only: it must not insert a new atom at the new rank zero. Compression adds existing cell masses, including atoms; it does not reconstruct a continuous density by dividing by cell volume. The hypotheses of Theorem 8.1 therefore apply. Alternatively, one can transform the generators by (15) and call an ordinary HV solver. Neither route improves the asymptotic exponent.

For Example A.5, the axis coordinates are $(0, 1, 2)$. The HV masses are $(0, 1, 1)$ and the magnitude masses are $(1, 1/2, 1/2)$. The only uncovered cell in the enclosing rank grid is $(2, 2)$, giving HV $4 - 1 = 3$ and magnitude $4 - 1/4 = 15/4$.

The public interface also supports the indicator directly:

```
from NumericalChanHVND.numerical_chan_local_compiled import magnitude
points = [(2, 1, 1, 1), (1, 2, 1, 1)]
assert abs(magnitude(points) - 135.0 / 16.0) < 1e-10
```

Here $(15/4)(3/2)^2 = 135/16$. An empty input has magnitude zero, whereas a single all-zero generator has magnitude one. Zero-extent boxes must therefore be retained for this application. The reflection from a minimization region to upper-corner extents is an ℓ_1 isometry, so it also preserves the metric magnitude of the corresponding region.

The saved validation record includes 18 checks of the two HV–magnitude transformations. The 64-point checkpoint instance and both checkpoint decisions were also checked for magnitude against transformed Chan $d/2$ inputs. The 1,552-box instance has magnitude 411 863 688.75, again agreeing with the transformed reference. The shared algebra tests include boundary atoms. These are correctness checks, not evidence that magnitude improves the running time of ordinary HV. The larger benchmark in Section 10 measures ordinary hypervolume only.

B A compact implementation and proof map

The following table locates the main mathematical obligations. The names refer to files in `NumericalChanHVND`; the proof specifies exact arithmetic while the compiled code evaluates the same formulas in binary64.

Mathematical operation	Implementation location and invariant
Axis cells and masses	<code>numerical_chan_numba.py</code> , <code>compute</code> : interval lengths for HV; ranks preserve box membership.
Cell simplification	<code>numerical_chan_compiled.py</code> , <code>_cell_state</code> : disjointness, slab tightening, cropping and pair masks.
Scalar contraction	<code>compiled_terms.py</code> , <code>eliminate</code> : first-winner guards and signed prefix differences.
Block compression	<code>compiled_terms.py</code> , <code>push_blocks</code> : lift to $2d$ variables, eliminate old axes, retain block masses.
Local schedule	<code>numerical_chan_local_compiled.py</code> , <code>_node</code> and <code>_new_block</code> : one budget step per cyclic advance; restart at every checkpoint.
Validation	<code>verify_compiled.py</code> checks the shared algebra; <code>verify_local_compiled.py</code> adds local scheduling and end-to-end checks.

The larger experiment uses `benchmarks/larger_4d_6d`, with saved coordinates, every retained repetition, partial runs and source hashes. Its unchanged solver sources are from revision `095fa73`. The earlier experiment remains in `benchmarks/compiled_2d_10d`, including its exact-oracle record and natural-compression probe. Revision `38587e7` identifies that earlier release. The local driver is now timed explicitly in the larger comparison; old speedup ratios are not transferred to it.

To reproduce correctness and build this manuscript from the repository root, use the following commands. LaTeX is run in the paper directory because the experimental sections and magnitude appendix are companion source files.

```
python NumericalChanHVND/verify_local_compiled.py
cd NumericalChanHVND/paper
pdflatex -interaction=nonstopmode -halt-on-error numerical_product_chan.tex
pdflatex -interaction=nonstopmode -halt-on-error numerical_product_chan.tex
```

The report uses the standard article class, A4 paper, 12-point text, and an inline `thebibliography` with `bibitem` entries. It can be compiled without a bibliography database or code execution by LaTeX.

References

- [1] Timothy M. Chan. Klee’s Measure Problem Made Easy. In *Proceedings of the 54th Annual IEEE Symposium on Foundations of Computer Science (FOCS)*, pages 410–419, 2013. <https://doi.org/10.1109/FOCS.2013.51>. Author version, 10 August 2013: https://tmc.web.engr.illinois.edu/easyklee8_13.pdf.
- [2] Michael T. M. Emmerich. The Magnitude of Dominated Sets: A Pareto Compliant Indicator Grounded in Metric Geometry. arXiv:2604.18147v1, 2026. <https://doi.org/10.48550/arXiv.2604.18147>.
- [3] Tom Leinster and Mark W. Meckes. The magnitude of a metric space: from category theory to geometric measure theory. In Nicola Gigli, editor, *Measure Theory in Non-Smooth Spaces*, pages 156–193. De Gruyter, 2017. <https://doi.org/10.1515/9783110550832-005>. Preprint: <https://arxiv.org/abs/1606.00095>.
- [4] Michael T. M. Emmerich and André H. Deutz. A tutorial on multiobjective optimization: fundamentals and evolutionary methods. *Natural Computing*, 17:585–609, 2018. <https://doi.org/10.1007/s11047-018-9685-y>.
- [5] Michael T. M. Emmerich. Pair-Elimination and Product-Measure Simplification of Chan’s Fixed-Dimensional Hypervolume Algorithm. Working research note, 2 October 2026, in the HV4DMagnitude repository. [Report PDF](#); [LaTeX source](#). <https://github.com/emmerichmtm/HV4DMagnitude>.
- [6] *FastHVChan: numerical hypervolume and magnitude implementations*. Repository of Michael Emmerich, including `NumericalChanHV4D` and `NumericalChanHVND`. Solver releases: `38587e7` (archived small-input benchmark) and `095fa73` (larger 4D–6D comparison), 4 October 2026. Complete source hashes and samples: accompanying `NumericalChanHVND` benchmark files. <https://github.com/emmerichmtm/FastHVChan>.